

Lab report

Inheritance

In order to explore the assembly language that is created with classes, I created a simple “Numbers” class which holds three integer pointers, and a subclass called “MoreNumbers” which holds an additional integer pointer. The code for the class Numbers is shown in Figure 1 and the code for the subclass MoreNumbers is shown in Figure 2.

```
class Numbers {
public:
    Numbers( int *n1, int *n2, int *n3 );
    ~Numbers();
    int *x1;
    int *x2;
    int *x3;
};

Numbers::Numbers( int *n1, int *n2, int *n3 ) {
    x1 = n1;
    x2 = n2;
    x3 = n3;
}

Numbers::~~Numbers() {
    delete x1;
    delete x2;
    delete x3;
}
```

Figure 1

```
class MoreNumbers : public Numbers {
public:
    MoreNumbers( int *n1, int *n2, int *n3, int *n4 );
    ~MoreNumbers();
    int *x4;
};

MoreNumbers::MoreNumbers( int *n1, int *n2, int *n3, int *n4 ) : Numbers ( n1, n2, n3 ) {
    x4 = n4;
}

MoreNumbers::~~MoreNumbers() {
    delete x4;
}
```

Figure 2

As we can see in Figure 3 the main function of the program simply instantiates various integers, and creates one instance of Numbers and one of MoreNumbers, and then deletes them in order to be able to properly examine the constructor and destructor of both a normal object and an object which inherits some of its data.

```
int main() {
    int t1 = 0;
    int t2 = 2;
    int t3 = 3;
    Numbers *n = new Numbers ( &t1, &t2, &t3 );

    t1 = 5;
    n->x1 = &t1;

    t1 = 9;
    t2 = 10;
    t3 = 16;
    int t4 = 15;

    MoreNumbers *m = new MoreNumbers( &t1, &t2, &t3, &t4 );

    delete n;
    delete m;
}
```

Figure 3

The first few lines in the snippet are showing the three arguments for the Numbers constructor (0, 2, 3) as written in the C++. Below the lines of argument we can see the command move eax, 24. From the online search, I found that it may be passing in the size of the arguments ($24 = 8 * 3$, 8 for each pointer). The function that is called is called Znmw doesn't seem to be in any part of the assembly file, and it is called again later in the code to call the constructor for MoreNumbers, I think this is probably to do the inherited functionality from the base-class's constructor. However, it appears to be a function that is called before the constructor, which I believe to be the function ZN7NumbersC1EPiS0_S0_. The next thing I noticed was that it appears that the return value (the object itself) is stored 40 bytes behind the base pointer. I know that 44 bytes behind the base pointer lies the t1 member variable, as command mov dword ptr [rbp - 8], 5 seems to be setting it to five. This means that the first four bytes (40-44 behind the base) are probably the implicit "this", and the following four byte chunks are the data members.

```
.Ltmp29:
    .cfi_def_cfa_register rbp
    sub    rsp, 112
    mov     dword ptr [rbp - 4], 0
    mov     dword ptr [rbp - 8], 0
    mov     dword ptr [rbp - 12], 2
    mov     dword ptr [rbp - 16], 3
    mov     eax, 24
    mov     edi, eax
    call    _Znmw
    mov     rdi, rax
    mov     rcx, rax

.Ltmp15:
    lea     rsi, [rbp - 8]
    lea     rdx, [rbp - 12]
    lea     r8, [rbp - 16]
    mov     qword ptr [rbp - 56], rdi # 8-byte Spill
    mov     rdi, rax
    mov     qword ptr [rbp - 64], rcx # 8-byte Spill
    mov     rcx, r8
    call    _ZN7NumbersC1EPiS0_S0_

.Ltmp16:
    jmp     .LBB5_1

.LBB5_1:
    mov     rax, qword ptr [rbp - 64] # 8-byte Reload
    mov     qword ptr [rbp - 24], rax
    mov     dword ptr [rbp - 8], 5
    mov     rcx, qword ptr [rbp - 24]
    lea     rdx, [rbp - 8]
    mov     qword ptr [rcx], rdx
    mov     dword ptr [rbp - 8], 9
    mov     dword ptr [rbp - 12], 10
    mov     dword ptr [rbp - 16], 16
    mov     dword ptr [rbp - 40], 15
    mov     esi, 32
```

```
    mov     esi, 32
    mov     edi, esi
    mov     qword ptr [rbp - 72], rdx # 8-byte Spill
    call    _Znmw
    mov     rcx, rax
    mov     rdx, rax

.Ltmp18:
    lea     rdi, [rbp - 12]
    lea     r8, [rbp - 16]
    lea     r9, [rbp - 40]
    mov     qword ptr [rbp - 80], rdi # 8-byte Spill
    mov     rdi, rax
    mov     rsi, qword ptr [rbp - 72] # 8-byte Reload
    mov     rax, qword ptr [rbp - 80] # 8-byte Reload
    mov     qword ptr [rbp - 88], rdx # 8-byte Spill
    mov     rdx, rax
    mov     qword ptr [rbp - 96], rcx # 8-byte Spill
    mov     rcx, r8
    mov     r8, r9
    call    _ZN11MoreNumbersC1EPiS0_S0_S0_

.Ltmp19:
    jmp     .LBB5_2

.LBB5_2:
    mov     rax, qword ptr [rbp - 88] # 8-byte Reload
    mov     qword ptr [rbp - 48], rax
    mov     rcx, qword ptr [rbp - 24]
    cmp     rcx, 0
    mov     qword ptr [rbp - 104], rcx # 8-byte Spill
    .LBB5_5
    je

# BB#3:
.Ltmp21:
    mov     rdi, qword ptr [rbp - 104] # 8-byte Reload
    call    _ZN7NumbersD1Ev

.Ltmp22:
    jmp     .LBB5_4

.LBB5_4:
    mov     rax, qword ptr [rbp - 104] # 8-byte Reload
    mov     rdi, rax
    call    _ZdlPv
```

Figure 4

Command line `ZN11NumbersC1EPisS0_S0_ S0_` seems to be the call to the constructor of `MoreNumbers`, and in name alone, it seems to have a lot in common with the constructor of `Numbers`, except for the addition of “More” and one extra “S0_” suffix. This is probably a result of the inheritance. I was expecting the assembler to automatically call both constructors, perhaps calling the base-class’s constructor as a part of the call to the constructor of the subclass, and I predicted correctly.

```
.text
.intel_syntax noprefix
.file "classOptimize.cpp"
.section .text.startup,"ax",@progbits
.align 16, 0x90
.type __cxx_global_var_init,@function
__cxx_global_var_init:                # @__cxx_global_var_init
.cfi_startproc
# BB#0:
push    rbp
.Ltmp0:
.cfi_def_cfa_offset 16
.Ltmp1:
.cfi_offset rbp, -16
mov     rbp, rsp
.Ltmp2:
.cfi_def_cfa_register rbp
sub     rsp, 16
movabs  rdi, ZStL8_ioinit
call    _ZNSt8ios_base4InitC1Ev
movabs  rdi, _ZNSt8ios_base4InitD1Ev
movabs  rsi, ZStL8_ioinit
movabs  rdx, __dso_handle
call    __cxa_atexit
mov     dword ptr [rbp - 4], eax # 4-byte Spill
add     rsp, 16
pop     rbp
ret
.Lfunc_end0:
```

Figure 5

```
.globl _ZN11MoreNumbersC2EPisS0_S0_
.align 16, 0x90
.type _ZN11MoreNumbersC2EPisS0_S0_,@function
_ZN11MoreNumbersC2EPisS0_S0_:        # @_ZN11MoreNumbersC2EPisS0_S0_
.cfi_startproc
# BB#0:
push    rbp
.Ltmp9:
.cfi_def_cfa_offset 16
.Ltmp10:
.cfi_offset rbp, -16
mov     rbp, rsp
.Ltmp11:
.cfi_def_cfa_register rbp
sub     rsp, 64
mov     qword ptr [rbp - 8], rdi
mov     qword ptr [rbp - 16], rsi
mov     qword ptr [rbp - 24], rdx
mov     qword ptr [rbp - 32], rcx
mov     qword ptr [rbp - 40], r8
mov     rcx, qword ptr [rbp - 8]
mov     rdx, rcx
mov     rsi, qword ptr [rbp - 16]
mov     rdi, qword ptr [rbp - 24]
mov     r8, qword ptr [rbp - 32]
mov     qword ptr [rbp - 48], rdi # 8-byte Spill
mov     rdi, rdx
mov     rdx, qword ptr [rbp - 48] # 8-byte Reload
mov     qword ptr [rbp - 56], rcx # 8-byte Spill
mov     rcx, r8
call    _ZN7NumbersC2EPisS0_S0_
mov     rcx, qword ptr [rbp - 40]
mov     rdx, qword ptr [rbp - 56] # 8-byte Reload
mov     qword ptr [rdx + 24], rcx
add     rsp, 64
pop     rbp
ret
```

Figure 6

Figure 5 shows the body of the `Numbers` constructor, and Figure 6 shows the body of the `MoreNumbers` constructor. As is notorious on Figure 6, we can see the command `ZN7NumbersC2EPisS0_S0_` in the `MoreNumbers` constructor calls the `Numbers` constructor within itself. However, I could not exactly tell how the data is manipulated after the constructor is called though.

I believe that per the x86 naming conventions of clang++, the constructors are always the name of the class followed by a C, and the destructors with a D. The destructor functions are the next called in the main function, as we see in Figure 7 the two command lines are:

`_ZN7NumbersD1Ev` and `_ZN11NumbersD1Ev`.

```

        mov     qword ptr [rbp - 104], rcx # 8-byte Spill
        je      .LBB5_5
# BB#3:
.Ltmp21:
        mov     rdi, qword ptr [rbp - 104] # 8-byte Reload
        call    _ZN7NumbersD1Ev
.Ltmp22:
        jmp     .LBB5_4
.LBB5_4:
        mov     rax, qword ptr [rbp - 104] # 8-byte Reload
        mov     rdi, rax
        call    _ZdlPv
.LBB5_5:
        mov     rax, qword ptr [rbp - 48]
        cmp     rax, 0
        mov     qword ptr [rbp - 112], rax # 8-byte Spill
        je      .LBB5_8
# BB#6:
.Ltmp24:
        mov     rdi, qword ptr [rbp - 112] # 8-byte Reload
        call    _ZN11MoreNumbersD1Ev
.Ltmp25:
        jmp     .LBB5_7
.LBB5_7:
        mov     rax, qword ptr [rbp - 112] # 8-byte Reload
        mov     rdi, rax
        call    _ZdlPv
.LBB5_8:
        mov     eax, dword ptr [rbp - 4]
        add     rsp, 112
        pop     rbp
        ret
.LBB5_9:

```

Figure 7

The prologue for the first destructor seems to be data passed in through the RDI register from behind the base pointer. This is probably the “this” pointer for the object that is allocated somewhere on the heap. The next function called with “_ZdlPv” is kind of confusing, but from purely the name it sounds like it deletes a pointer variable, and perhaps it is responsible for deallocating the data members of the instance of the Numbers class. From searching through the Internet, I found references to this strange function on gnu.org and realized that it is a part of GCC and stands for “operator delete (void*)” meaning it deletes a void pointer (a point of any type). It seems to me like this is probably the function used to delete all pointers in every single destructor.

As we can see in the picture, the function is called again after the MoreNumbers destructor is called, this makes me believe that this Zdlpv is the function to actually free the memory itself, whereas the other functions that share the namesake of their respective classes exist simply to perform any other functionality that might need to exist for an objects destruction.

Dynamic Dispatch

To explore how Dynamic dispatch works in Assembly code I created a program where I had to use virtual functions, as we can see in Figure 1, to help the compiler determine the runtime of the function. The line `x.f()` gets executed twice, but a different function gets called each time. Dynamic dispatch is implemented by means of a virtual function table that has the address of the final overrider for the complement object.

```
#include <iostream>

class base {
public:
    virtual void f() const = 0 ;
    virtual ~base() {}
};

class A : public base {
public:
    virtual void f() const { std::cout << "A::f()" << std::endl; }
};

class B : public base {
public:
    virtual void f() const { std::cout << "B::f()" << std::endl; }
};

void dispatch(const base & x) {
    x.f();
}

int main() {
    A a ;
    B b ;

    dispatch(a) ;
    dispatch(b) ;
}
```

Figure 1

After creating the Assembly code the first thing I noticed was the opcode `movabs`, which was something I have not seen before in any of my programs, was being used for the passing values of the function `_ZNSt8ios_base4InitC1Ev`. The searching about the meaning of this opcode I found that is just a GAS specific way to enforce encoding in a 64 bit memory offset. `Movabs` should perform the same way as the standard move opcode.

Figure 2 and Figure 3 show how the line `x.f()` gets called twice in the dispatch function. The lines where virtual dispatch takes place in Assembly are shown in Figure 3. From the picture we can see that the line `mov rdi, qword ptr [rdi - 8]` is loading the vtable from the object, then loading the function address from the vtable to the `rax` register in the following line: `mov rax, qword ptr [rdi]` (`rdi` represents the value that 'this' is pointing to), then calling the function directly (`call qword ptr [rax]`). So we see that the address is fetched and called all in one statement. Finally the stack pointer is moved back up to clean off the arguments that were

pushed before the call. Depending on compiler implementations, in C++ is common to see the caller clean off the arguments.

```
.text
.intel_syntax noprefix
.file "base2.cpp"
.section .text.startup,"ax",@progbits
.align 16, 0x90
.type __cxx_global_var_init,@function
__cxx_global_var_init:                # @__cxx_global_var_init
.cfi_startproc
# BB#0:
push    rbp
.Ltmp0:
.cfi_def_cfa_offset 16
.Ltmp1:
.cfi_offset rbp, -16
mov     rbp, rsp
.Ltmp2:
.cfi_def_cfa_register rbp
sub     rsp, 16
movabs  rdi, _ZNStL8_ioinit
call    _ZNSt8ios_base4InitC1Ev
movabs  rdi, _ZNSt8ios_base4InitD1Ev
movabs  rsi, _ZNStL8_ioinit
movabs  rdx, __dso_handle
call    __cxa_atexit
mov     dword ptr [rbp - 4], eax # 4-byte Spill
add     rsp, 16
pop     rbp
.Lfunc_end0:
.size   __cxx_global_var_init, .Lfunc_end0-__cxx_global_var_init
.cfi_endproc

.text
.globl _Z8dispatchRK4base
.align 16, 0x90
.type _Z8dispatchRK4base,@function
```

Figure 2

```
_Z8dispatchRK4base:                # @_Z8dispatchRK4base
.cfi_startproc
# BB#0:
push    rbp
.Ltmp3:
.cfi_def_cfa_offset 16
.Ltmp4:
.cfi_offset rbp, -16
mov     rbp, rsp
.Ltmp5:
.cfi_def_cfa_register rbp
sub     rsp, 16
mov     qword ptr [rbp - 8], rdi
mov     rdi, qword ptr [rbp - 8]
mov     rax, qword ptr [rdi]
call    qword ptr [rax]
add     rsp, 16
pop     rbp
.Lfunc_end1:
.size   _Z8dispatchRK4base, .Lfunc_end1-_Z8dispatchRK4base
.cfi_endproc

.globl main
.align 16, 0x90
.type   main,@function
```

Figure 3

The code below shows the part of the main method. The end of the dispatch function ends with “sub rsp, 48” and the epilog starts with the instruction “add rsp, 48”. As we see in the picture, the instructions between the prolog and epilog access stack content using RSP as a reference without having any push and pop instructions intervening in the function body.

```

.cfi_offset rbp, -16
mov     rbp, rsp
.Ltmp21:
.cfi_def_cfa_register rbp
sub     rsp, 48
lea     rax, [rbp - 8]
mov     rdi, rax
mov     qword ptr [rbp - 40], rax # 8-byte Spill
call    _ZN1AC2Ev
lea     rdi, [rbp - 16]
call    _ZN1BC2Ev

.Ltmp6:
mov     rdi, qword ptr [rbp - 40] # 8-byte Reload
call    _Z8dispatchRK4base

.Ltmp7:
jmp     .LBB2_1

.LBB2_1:
.Ltmp8:
lea     rdi, [rbp - 16]
call    _Z8dispatchRK4base

.Ltmp9:
jmp     .LBB2_2

.LBB2_2:
.Ltmp13:
lea     rdi, [rbp - 16]
call    _ZN1BD2Ev

.Ltmp14:
jmp     .LBB2_3

.LBB2_3:
lea     rdi, [rbp - 8]
call    _ZN1AD2Ev
xor     eax, eax
add     rsp, 48

```

Figure 4

Sources

<http://lists.gnu.org/archive/html/bugbinutils/2010-01/msg00047.html>

http://www.codemachine.com/article_x64deepdive.html

<http://www.stackoverflow.com/questions/20147054/hot-does-dynamic-dispatch-happen-in-assembly>